

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>
(<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[7.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
#importing libraries
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```
C:\Users\HIMANSHU NEGI\Anaconda3\lib\site-packages\gensim\utils.py:1212: U
serWarning: detected Windows; aliasing chunkize to chunkize_serial
  warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
```

[1]. Reading Data

In [2]:

```
# using the SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

#filtering only positive and negative reviews i.e.
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5000 """, con)

# Give reviews with Score>3 a positive rating, and reviews with a score<3 a negative rating.
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (5000, 10)

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

In [3]:

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

In [4]:

```
print(display.shape)
display.head()
```

(80668, 7)

Out[4]:

	UserId	ProductId	ProfileName	Time	Score	Text
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...

In [5]:

```
display[display['UserId']== 'AZY10LLTJ71NX']
```

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommend to try green tea extract to ...

In [6]:

```
display['COUNT(*)'].sum()
```

Out[6]:

393063

Exploratory Data Analysis

[2] Data Cleaning: Deduplication


It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it is necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [7]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Help
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2



As can be seen above the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [8]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

In [9]:

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

Out[9]:

(4986, 10)

In [10]:

```
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[10]:

99.72

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [11]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[11]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Help
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

In [12]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [13]:

```
#Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

(4986, 10)

Out[13]:

```
1    4178
0     808
Name: Score, dtype: int64
```

In [14]:

```
final['Text'].shape
```

Out[14]:

(4986,)

[3]. Text Preprocessing.

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [15]:

```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

Why is this \$[...] when the same product is available for \$[...] here?
<http://www.amazon.com/VICTOR-FLY-MAGNET-BAIT-REFILL/dp/B00004RBDY>

The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there were a lot of "brown" chips in the bsg (my favorite), so I bought some more through amazon and shared with my family and friends. I am a little disappointed that there are not, so far, very many brown chips in these bags, but the flavor is still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salty, and the onion flavor is better. If you haven't eaten Kettle chips before, I recommend that you try a bag before buying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

=====

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering.

These are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion.

Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blur would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.

So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

=====

love to order my coffee on amazon. easy and shows up quickly.
This cup is great coffee. dcaf is very good as well

=====

In [16]:

```
# remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

In [17]:

```
# https://stackoverflow.com/questions/16206380/python-beautifulsoup
# removing-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1001, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1502, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4903, 'lxml')
text = soup.get_text()
print(text)
```

Why is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

=====

We are always pleased with Amazon's packaging. Chips are never crushed or broken, always fresh. We really liked this brand/flavor of chips. Very tasty!

=====

I purchased these thinking they would just be ordinary oatmeal raisin cookies. Wrong! I like chocolate chip and I like oatmeal raisin but these combination cookies just don't work well together. I wish I had read the item listing more carefully. Even my kids said "Blech!".

=====

This is a great coffee. Its dark but not bitter, a little chocolaty, and it reminds me of an Adirondack lodge. I can get them for a better price in a Caribou Coffee store, however.

In [18]:

```
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"'re", " are", phrase)
    phrase = re.sub(r"'s", " is", phrase)
    phrase = re.sub(r"'d", " would", phrase)
    phrase = re.sub(r"'ll", " will", phrase)
    phrase = re.sub(r"'t", " not", phrase)
    phrase = re.sub(r"'ve", " have", phrase)
    phrase = re.sub(r"'m", " am", phrase)
    return phrase
```

In [19]:

```
sent_1502 = decontracted(sent_1502)
print(sent_1502)
print("="*50)
```

I purchased these thinking they would just be ordinary oatmeal raisin cookies. Wrong! I like chocolate chip and I like oatmeal raisin but these combination cookies just do not work well together. I wish I had read the item listing more carefully. Even my kids said "Blech!".

=====

In [20]:

```
#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />
The Victor and traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

In [21]:

```
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1502 = re.sub('[^A-Za-z0-9]+', ' ', sent_1502)
print(sent_1502)
```

I purchased these thinking they would just be ordinary oatmeal raisin cookies Wrong I like chocolate chip and I like oatmeal raisin but these combination cookies just do not work well together I wish I had read the item listing more carefully Even my kids said Blech

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourself', 's', 'you', 'you're", "you've", \
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 't hey', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "th at'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'ha d', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as' , 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through' , 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'ov er', 'under', 'again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'an y', 'both', 'each', 'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too' , 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no w', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'migh tn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'w asn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"]])
```

```
# Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())
```

```
100% |██████████████████████████████████████████████████████████████████████|  
██████████ | 4986/4986 [00:02<00:00, 1853.76it/s]
```

In [24]:

```
preprocessed_reviews[1502]
```

Out[24]:

```
'purchased thinking would ordinary oatmeal raisin cookies wrong like choco  
late chip like oatmeal raisin combination cookies not work well together w  
ish read item listing carefully even kids said blech'
```

[3.2] Preprocess Summary for reviews

In [93]:

```
## preprocessing for review summary.  
final['Summary'].shape  
print(final['Summary'])
```

2546 thirty bucks?
 2547 Flies Begone
 1145 WOW Make your own 'slickers' !
 1146 Great Product
 2942 Good stuff!
 2941 Premium Quality Dog Food!!!
 1071 Cats love it!
 2187 Nice, Big Pieces & Big Almond Flavor
 4695 A Summer Treat Fat Free, Guilt Free
 2068 Don't buy this product unless you are looking ...
 2069 Little Flavor
 2806 A Staple in my house
 2805 A favorite quick meal solution
 4099 Best Hot Sauce and Taco Sauce Available In Ame...
 4096 Pico Pica ... the BEST
 4097 This is the stuff!
 4098 What everyone is saying here about Pico Pica i...
 1332 not edible
 1330 shining star
 1329 The Inexpensive Alternative to Gold Leaf!
 1328 Create Exquisite Cake Decorations
 1331 gold dust is awesome
 4320 Perfect for my sons cake.
 4321 Really cute - made a great golf cake
 4322 Golf Set
 4323 Cake Topper
 4054 Fantastic
 2477 Adzuki beans
 2476 Yum.
 2478 yummy
 ...
 2214 burnt toast
 2215 Love Caribou!
 2212 down to almost none
 2216 yum
 2217 great coffee
 677 cute ,cute, cute!
 678 A Surprising Find
 3663 It's the best
 3662 Huge Fan
 3664 Awesome sauce!
 3580 Good Malta, baffling business model
 1110 Service was good
 1109 I was able to eat bread again!
 1108 Love Taro
 1107 Delicious and easy
 1106 Favorite thing about Brazil
 1232 Delicious
 4714 Best Italian olive oil
 2013 Mild Taste, But Delicious.
 3567 Great drink, horrible price!
 3271 More for Me
 220 OMG best chocolate jelly belly
 4117 GREAT STUFF
 4118 You won't believe it
 712 One of the better T-Discs
 711 great coffee - terrible price
 710 Best of the Tassimo's
 709 Good Tasting cup o' joe
 713 Kona for Tassimo

1362 weak coffee not good for a premium product and...
 Name: Summary, Length: 4986, dtype: object

In [158]:

```
# printing some random reviews
summ_0 = final['Summary'].values[0]
print(summ_0)
print("="*50)

summ_1000 = final['Summary'].values[1000]
print(summ_1000)
print("="*50)

summ_1500 = final['Summary'].values[1500]
print(summ_1500)
print("="*50)

summ_4900 = final['Summary'].values[4900]
print(summ_4900)
print("="*50)
```

```
thirty bucks?
=====
Best sour cream & onion chip I've had
=====
Are We Reviewing Our Mistakes Or These Cookies?
=====
caribou
=====
```

In [159]:

```
# remove urls from text python:
summ_0 = re.sub(r"http\S+", "", summ_0)
summ_1000 = re.sub(r"http\S+", "", summ_1000)
summ_1500 = re.sub(r"http\S+", "", summ_1500)
summ_4900 = re.sub(r"http\S+", "", summ_4900)

print(summ_0)
```

```
thirty bucks?
```

In [160]:

```
# remove urls from text python:
i=0;
for sm in x:
    if (len(re.findall('<.*?>', sm))):
        summ = re.sub(r"http\S+", "", sm)
        print(i)
        print(sent)

    i += 1;

print(sm)
```

```
weak coffee not good for a premium product and price
```

In [161]:

```
# removing-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(summ_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(summ_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(summ_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(summ_4900, 'lxml')
text = soup.get_text()
print(text)
```

```
thirty bucks?
=====
Best sour cream & onion chip I've had
=====
Are We Reviewing Our Mistakes Or These Cookies?
=====
caribou
```

In [151]:

```
# function to clean the word of any punctuation or special characters
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

In [162]:

```
summ_1500 = decontracted(summ_1500)
print(summ_1500)
print("="*50)
```

Are We Reviewing Our Mistakes Or These Cookies?

=====

In [163]:

```
#remove words with numbers python:
summ_0 = re.sub("\S*\d\S*", "", summ_0).strip()
print(summ_0)
```

thirty bucks?

In [164]:

```
#remove spacial character:
summ_1500 = re.sub('[^A-Za-z0-9]+', ' ', summ_1500)
print(summ_1500)
```

Are We Reviewing Our Mistakes Or These Cookies

In [155]:

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourself', 't
s', 'you', "you're", "you've", \
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'his', 'himself', \
    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 't
hey', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "th
at'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'ha
d', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as'
, 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through'
, 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'ov
er', 'under', 'again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'an
y', 'both', 'each', 'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too'
, 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no
w', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
'doesn', "doesn't", 'hadn', \
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'migh
tn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'w
asn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

In [168]:

```
# Combining all the above stundents
from tqdm import tqdm
preprocessed_summary = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Summary'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwor
ds)
    preprocessed_summary.append(sentence.strip())
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 4986/4986 [00:01<00:00, 2627.73it/s]
```

In [169]:

```
preprocessed_summary[1500]
```

Out[169]:

```
'reviewing mistakes cookies'
```

[4] Featurization

[4.1] BAG OF WORDS

In [25]:

```
#Bow
count_vect = CountVectorizer() #function in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])

some feature names ['aa', 'aahhs', 'aback', 'abandon', 'abates', 'abbot
t', 'abby', 'abdominal', 'abiding', 'ability']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 12997)
the number of unique words 12997
```

[4.2] Bi-Grams and n-Grams.

In [26]:

```
#bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html
# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_c
ounts.get_shape()[1])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

[4.3] TF-IDF

In [27]:

```
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names
()[0:10])
print('='*50)

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.g
et_shape()[1])

some sample features(unique words in the corpus) ['ability', 'able', 'able
find', 'able get', 'absolute', 'absolutely', 'absolutely delicious', 'abso
lutely love', 'absolutely no', 'according']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

In [32]:

```
# average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to
    change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 4986/4986 [00:07<00:00, 706.58it/s]
```

```
4986
50
```

In [33]:

In [34]:

```
100%|███████████████████████████████████████████████████████████████████████████|
██████████ | 4986/4986 [00:33<00:00, 147.21it/s]
```

[5.0] Applying TNSE on Text BOW vectors

In [36]:

http://localhost:8888/nbconvert/html/amazon%20fine%20food/02%20Amazon%20Fine%20Food%20Reviews%20Analysis TSNE.ipynb?downl... 23/32

In [63]:

```
type(final_counts)
```

Out[63]:

```
numpy.matrixlib.defmatrix.matrix
```

In [41]:

```
print(final_counts.shape)
```

```
(4986, 12997)
```

In [52]:

```
final['Score'].value_counts()
```

Out[52]:

```
1    4178
```

```
0     808
```

```
Name: Score, dtype: int64
```

In [62]:

```
print(final['Score'].shape)
```

```
(4986,)
```

t-SNE of Bag of Words(BoW) with perplexity = 30 and n_iter = 1000


In [66]:

```
from sklearn.manifold import TSNE
model = TSNE(n_components=2, random_state=0, perplexity = 30, n_iter = 1000)
# configuring the parameteres
# the number of components = 2
# default perplexity = 30
# default learning rate = 200
# default Maximum number of iterations for the optimization = 1000

tsne_dataBow = model.fit_transform(final_counts)

# creating a new data frame which help us in plotting the result data
tsne_dataBow = np.vstack((tsne_dataBow.T, final['Score'])).T
tsnebow_df = pd.DataFrame(data=tsne_dataBow, columns=("dim1", "dim2", "score"))

# Ploting the result of tsne
sns.FacetGrid(tsnebow_df, hue="score", size=6).map(plt.scatter, 'dim1', 'dim2').add_
leg
end()
plt.title("TSNE for Bag of Words")
plt.show()
```



OBSERVATION:-

1> From above plots it is observed that around 90% and above points are overlapping so it is very difficult to classify the polarity of the reviews i.e positive or negative.

2> It is hard to classify points using a simple linear line or a plane.

3> Here are some points which are very far from most of the points which can be outliers and can easily affect the models

[5.1] Applying TNSE on Text TFIDF vectors

TFIDF vectors are sparse matrices so we have to convert it into dense matrices using `todense()` function

In [69]:

```
type(final_tf_idf)
```

Out[69]:

```
scipy.sparse.csr.csr_matrix
```

In [70]:

```
final_tf_idf = final_tf_idf.todense()
```

In [71]:

```
type(final_tf_idf)
```

Out[71]:

```
numpy.matrixlib.defmatrix.matrix
```

t-SNE of tf-idf with perplexity = 30 and n_iter = 1000

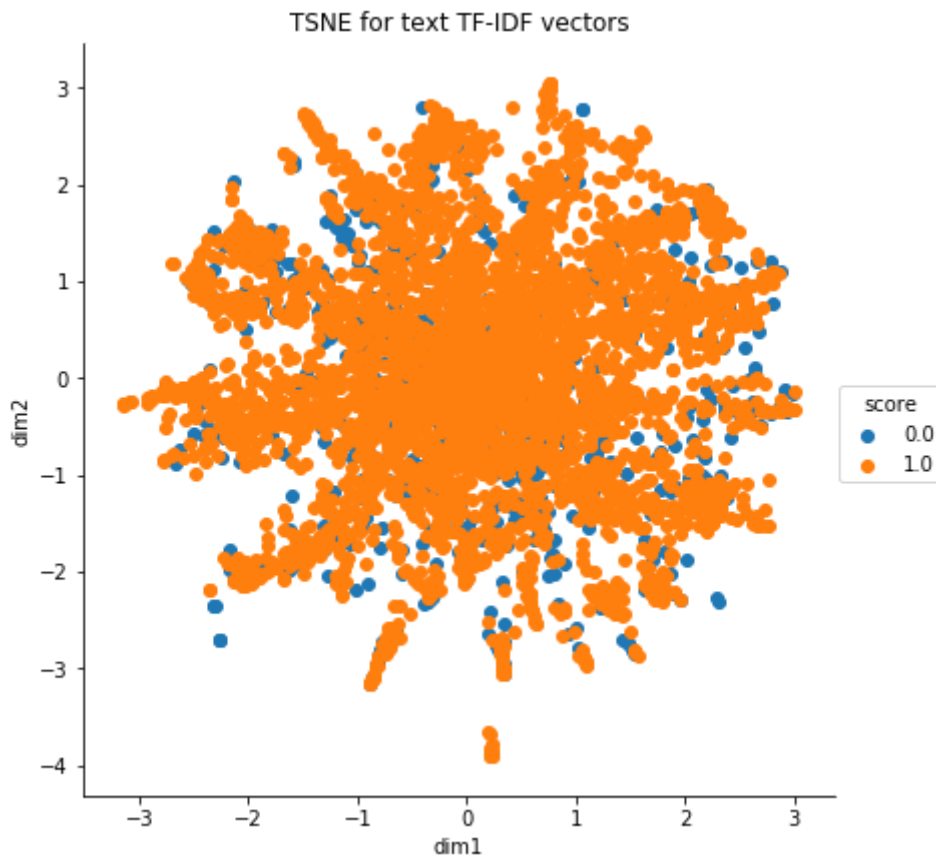
In [75]:

```
from sklearn.manifold import TSNE
model = TSNE(n_components=2, random_state=0, perplexity = 30, n_iter = 1000)
# configuring the parameters
# the number of components = 2
# default perplexity = 30
# default learning rate = 200
# default Maximum number of iterations for the optimization = 1000

tsne_datAtfidf = model.fit_transform(final_tf_idf)

# creating a new data frame which help us in plotting the result data
tsne_datAtfidf = np.vstack((tsne_datAtfidf.T, final['Score'])).T
tsnetfidf_df = pd.DataFrame(data=tsne_datAtfidf, columns=("dim1", "dim2", "score"))

# Ploting the result of tsne
sns.FacetGrid(tsnetfidf_df, hue="score", size=6).map(plt.scatter, 'dim1', 'dim2').add_
legend()
plt.title("TSNE for text TF-IDF vectors")
plt.show()
```



OBSERVATION:-

1> From above plots it is observed that around (85-90)% points are overlapping so it is very difficult to classify the polarity of the reviews i.e positive or negative.

2> It is hard to classify points using a simple linear model.

[5.2] Applying TNSE on Text Avg W2V vectors

t-SNE of avg W2V with perplexity = 30 and n_iter = 1000


In [74]:

```
from sklearn.manifold import TSNE
model = TSNE(n_components=2, random_state=0, perplexity = 30, n_iter = 1000)
# configuring the parameters
# the number of components = 2
# default perplexity = 30
# default learning rate = 200
# default Maximum number of iterations for the optimization = 1000

tsne_dataWv = model.fit_transform(sent_vectors)

# creating a new data frame which help us in plotting the result data
tsne_dataWv = np.vstack((tsne_dataWv.T, final['Score'])).T
tsneAwv_df = pd.DataFrame(data=tsne_dataWv, columns=("dim1", "dim2", "score"))

# Plotting the result of tsne
sns.FacetGrid(tsneAwv_df, hue="score", size=6).map(plt.scatter, 'dim1', 'dim2').add_
legend()
plt.title("TSNE for AVG Word2Vec")
plt.show()
```



OBSERVATION:-

1> From above plots it is observed that around 90% and above points are overlapping so it is very difficult to classify the polarity of the reviews i.e positive or negative.

2> It is hard to classify points using a simple linear line or a plane

[5.3] Applying TNSE on Text TFIDF weighted W2V vectors

t-SNE of tf-idf weighted W2V with perplexity = 30 and n_iter = 1000


In [76]:

```
from sklearn.manifold import TSNE
model = TSNE(n_components=2, random_state=0, perplexity = 30, n_iter = 1000)
# configuring the parameteres
# the number of components = 2
# default perplexity = 30
# default learning rate = 200
# default Maximum number of iterations for the optimization = 1000

tsne_datfidfw2v = model.fit_transform(tfidf_sent_vectors)

# creating a new data frame which help us in plotting the result data
tsne_datfidfw2v = np.vstack((tsne_datfidfw2v.T, final['Score'])).T
tsnefidfw2v_df = pd.DataFrame(data=tsne_datfidfw2v, columns=("dim1", "dim2", "score"))

# Ploting the result of tsne
sns.FacetGrid(tsneAwv_df, hue="score", size=6).map(plt.scatter, 'dim1', 'dim2').add_
leg
end()
plt.title("TSNE for TF-IDF weighted W2V vectors")
plt.show()
```



OBSERVATION:-

1> From above plots it is observed that around 90% and above points are overlapping so it is very difficult to classify the polarity of the reviews i.e positive or negative.

2> It is hard to classify points using a simple linear model.

[6] Conclusions

As we can see all above 4 t-SNE plots it is clearly shown that the data is highly overlapping due to which it is impossible to classify the positive and negative reviews/points with implementation a linear model or line or a plane.

As it is clear that our data is not linearly separable so we will make a model which will easily separate our data.

We have taken 5k points due to low system configuration if we will take more points then result can be better.

Due to low system configuration and t-SNE do lot of internal computations which take a lot of time and memory we have plot t-SNE on default perplexity = 30 ,default learning rate = 200,default number of iterations=1000 for better analysis we can plot it for more different parameters.

In []:

```
84.5, 83, 87, 86, 70.66.66, 73
```